Tu Vu

tuvu@vt.edu
https://tuvllms.github.io

Appointments

Virginia Tech   Fall 2024 — *present*
Assistant Professor, Computer Science
Research Interests: *natural language processing & machine learning*

Google
Faculty Researcher   Spring 2025 — *present*
Google DeepMind Research Scientist   2023 — 2024

Education

University of Massachusetts, Amherst   2016 — 2023
M.S/Ph.D. in Computer Science
Advisor: Mohit Iyyer   2018 — 2023

Thesis committee: Mohit Iyyer, Subhransu Maji, Hamed Zamani, Thang Luong, Colin Raffel

Vietnam National University, Hanoi   2009 — 2013
B.S. Honors Program in Computer Science
Highest distinction (class rank: 1/100)

Professional Experience

Google DeepMind   Fall 2022 — Spring 2023
Student Researcher *with* Thang Luong & Quoc Le

Google DeepMind   Summer 2022
Research Intern *with* Thibault Sellam & Elizabeth Clark

Google DeepMind   Winter 2021 — Spring 2022
Student Researcher *with* Noah Constant

Google DeepMind   Summer 2021 — Fall 2021
Research Intern & Student Researcher *with* Daniel Cer & Noah Constant

Google DeepMind   Winter 2020 — Spring 2021
Student Researcher *with* Thang Luong & Quoc Le

Google DeepMind   Summer 2020
Research Intern *with* Grady Simon & Zi Yang & Nan Hua

Microsoft Research   Summer 2019
Research Intern *with* Tong Wang & Tsendsuren Munkhdalai & Adam Trischler

Selected Preprints & Publications

*: equal contribution

For an up-to-date list of my research papers, please see my Google Scholar profile.

Foundational Autoraters: Taming Large Language Models for Better Automatic Evaluation

Tu Vu*, Kalpesh Krishna*, Salaheddin Alzubi, Chris Tar, Manaal Faruqui, Yun-Hsuan Sung

EMNLP 2024

// The top-performing generative model on RewardBench as of July 15, 2024, trained only on publicly available data

What Matters for Model Merging at Scale?

Prateek Yadav, **Tu Vu**, Jonathan Lai, Alexandra Chronopoulou, Manaal Faruqui, Mohit Bansal, Tsendsuren Munkhdalai

Under review @ ICLR 2025

Gemini: A Family of Highly Capable Multimodal Models

Google Gemini Team: Rohan Anil, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew Dai, Anja Hauth, and others including **Tu Vu**

arXiv preprint 2023

// Google AI Blog

FreshLLMs: Refreshing Large Language Models with Search Engine Augmentation

Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong

ACL 2024 Findings

// Our dataset and method have inspired or been used for the development of Google's Gemini, Perplexity.AI's Online LLMs, You.com, and Contextual AI's RAG 2.0

The Flan Collection: Designing Data and Methods for Effective Instruction Tuning

Shayne Longpre, Le Hou, **Tu Vu**, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc Le, Barret Zoph, Jason Wei, and Adam Roberts

ICML 2023

// Google Research Blog

Mixture-of-experts meets instruction tuning: A winning combination for large language models

Sheng Shen, Le Hou, Yanqi Zhou, Nan Du, Shayne Longpre, Jason Wei, Hyung Won Chung, Barret Zoph, William Fedus, Xinyun Chen, **Tu Vu**, Yuexin Wu, Wuyang Chen, Albert Webson, Yunxuan Li, Vincent Zhao, Hongkun Yu, Kurt Keutzer, Trevor Darrell, and Denny Zhou

ICLR 2024

SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer

Tu Vu, Brian Lester, Noah Constant, Rami Al-Rfou, and Daniel Cer

ACL 2022

Overcoming Catastrophic Forgetting in Zero-Shot Cross-Lingual Generation

Tu Vu, Aditya Barua, Brian Lester, Daniel Cer, Mohit Iyyer, and Noah Constant

EMNLP 2022

STraTA: Self-Training with Task Augmentation for Better Few-shot Learning

Tu Vu, Thang Luong, Quoc Le, Grady Simon, and Mohit Iyyer
EMNLP 2021

Exploring and Predicting Transferability across NLP Tasks
Tu Vu, Tong Wang, Tsendsuren Munkhdalai, Alessandro Sordoni, Adam Trischler, Andrew Mattarella-Micke, Subhransu Maji, and Mohit Iyyer
EMNLP 2020

FUNDING

ADVISING

PHD ADVISEES:

Quyet Do, 1st year PhD student at **Virginia Tech**   Fall 2024 — *present*

Thinh Pham, 1st year PhD student at **Virginia Tech**   Fall 2024 — *present*

Rishab Balasubramanian, 1st year PhD student at **Virginia Tech** Fall 2024 — *present*

Pin-Jie Lin, 1st year PhD student at **Virginia Tech**   Fall 2024 — *present*

OTHERS:

Prateek Yadav, Research Intern at **Google Gemini**   Summer 2024 — Spring 2025

Simeng Han, Student Researcher at **Google DeepMind**   Summer 2024 — Spring 2025

Salaheddin Alzubi, Masters student at **UMass Amherst**   Fall 2022 — Spring 2023

Dheeraj Mekala, PhD student at **UCSD**   Spring — Summer 2022

RECENT INVITED TALKS

Efficient Model Development in the Era of Large Language Models   November 2024
VinAI

Efficient Model Development in the Era of Large Language Models   October 2024
Mila / McGill NLP Seminar

Efficient Adaptation of Large Language Models   November 2023
Graph Neural Networks Reading Group, **Google**

Effective and Efficient Transfer Learning   Spring 2023
in the Era of Large Language Models
Faculty job talk

Overcoming Catastrophic Forgetting   October 2022
in Zero-Shot Cross-Lingual Generation
Parameter Efficient Tuning Methods Sync, **Google**

Transfer Learning with Large-scale Language Models   August 2022
Lecture at **The New Turing Institute**

The Appeal of Parameter-efficient Transfer Learning   June 2022
Natural Language Accelerated Team, **Google**

SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer   December 2021
Parameter Efficient Tuning Methods Sync, **Google**

ACADEMIC SERVICE

Area Chair for NAACL 2025, COLING 2025, ACL 2024, EMNLP 2024
Session Chair (Machine Learning for NLP) at EMNLP 2024
Program Committee/Reviewer for ICML 2025; NeurIPS 2024; TL4NLP@NeurIPS 2022; COLM 2024; AAAI 2023; ACL 2022, 2021, 2020, 2019; EMNLP 2022, 2021; NAACL 2022, 2021; COLING 2020; CoNLL 2019; INLG 2020, 2019

SELECTED MEDIA

Gemini: Google AI Blog   2023

FreshLLMs: ZDNET   2023

The Flan Collection: Google Research Blog   2023

SPoT: Headlines of Google AI's Natural Language Accelerated Newsletter   Q1, 2022

SELECTED AWARDS & HONORS

Google Student Researcherships   2020 — 2023

UMass Amherst Graduate Assistantships   2016 — 2023

Honda Y-E-S Award for young engineers and scientists, Vietnam   2013
// in the top 10 nationally

Outstanding Academic and Co-curricular Achievements, Vietnam National University   2013

Prominent Young Figure Award, Vietnam National University   2010 & 2012

First Runner-up Prize, International Programming Contest, Japan   2011
// ranked 2nd among 64 teams internationally

Outstanding Young Talent of the Capital City, Vietnam   2010
// in the top 100 most outstanding young talents selected from a wide range of fields

Champion Prize, National Mathematical Olympiad, Vietnam   2010
// ranked 1st among more than 600 contestants nationally

A number of prizes in National/International Olympiads   2009 — 2013
(in both Mathematics and Informatics)

A number of academic scholarships for undergraduate students   2009 — 2013

PATENTS

*: original inventor

Frozen Model Adaptation Through Soft Prompt Transfer
Tu Vu[*], Daniel Cer, Noah Constant, Brian Lester, Rami Al-Rfou
U.S. Patent Application, 17/863,840