

TU VU

 [tuvllms.github.io](https://github.com/tuvllms) |  tuvu@vt.edu |  [Google Scholar](#) |  [@tuvllms](#)

APPOINTMENTS

Virginia Tech

Fall 2024 — *present*

Assistant Professor, Computer Science

Research Interests: *large language models & natural language processing*

Google

Faculty Researcher

Spring 2025 — *present*

Google DeepMind Research Scientist

2023 — 2024

EDUCATION

University of Massachusetts, Amherst

2016 — 2023

M.S/PH.D. in Computer Science

Advisor: [Mohit Iyyer](#)

2018 — 2023

Thesis committee: [Mohit Iyyer](#), [Subhansu Maji](#), [Hamed Zamani](#), [Thang Luong](#), [Colin Raffel](#)

Vietnam National University, Hanoi

2009 — 2013

B.S. Honors Program in Computer Science

Highest distinction (class rank: 1st)

PROFESSIONAL EXPERIENCE

Google DeepMind

Fall 2022 — Spring 2023

Research Intern *with* [Thang Luong](#) & [Quoc Le](#)

Google DeepMind

Summer 2022

Research Intern *with* [Thibault Sellam](#) & [Elizabeth Clark](#)

Google DeepMind

Winter 2021 — Spring 2022

Research Intern *with* [Noah Constant](#)

Google DeepMind

Summer 2021 — Fall 2021

Research Intern *with* [Daniel Cer](#) & [Noah Constant](#)

Google DeepMind

Winter 2020 — Spring 2021

Research Intern *with* [Thang Luong](#) & [Quoc Le](#)

Google DeepMind

Summer 2020

Research Intern *with* [Grady Simon](#) & [Zi Yang](#) & [Nan Hua](#)

Microsoft Research

Summer 2019

Research Intern *with* [Tong Wang](#) & [Tsendsuren Munkhdalai](#) & [Adam Trischler](#)

*: equal contribution. For my latest research papers, see my [Google Scholar](#) profile.

SealQA: Raising the Bar for Reasoning in Search-Augmented Language Models

Thinh Pham, Nguyen Nguyen, Pratibha Zunjare, Weiyuan Chen, Yu-Min Tseng, **Tu Vu**

ICLR 2026

// Our benchmark dataset has been used by Google's Gemini, DeepSeek, and Kimi

Efficient Model Development through Fine-tuning Transfer

Pin-Jie Lin, Rishab Balasubramanian, Fengyuan Liu, Nikhil Kandpal, **Tu Vu**

EMNLP 2025

// Oral presentation, rating 4.5/5 by area chairs

What Matters for Model Merging at Scale?

Prateek Yadav, **Tu Vu**, Jonathan Lai, Alexandra Chronopoulou, Manaal Faruqui, Mohit Bansal, Tsendsuren Munkhdalai

TMLR 2025

Foundational Autoraters: Taming Large Language Models for Better Automatic Evaluation

Tu Vu*, Kalpesh Krishna*, Salaheddin Alzubi, Chris Tar, Manaal Faruqui, Yun-Hsuan Sung

EMNLP 2024

// The top-performing generative model on RewardBench as of July 15, 2024, trained only on publicly available data

Gemini: A Family of Highly Capable Multimodal Models

Google Gemini Team: Rohan Anil, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew Dai, Anja Hauth, and others including

Tu Vu

Technical report 2023

// Google AI Blog

FreshLLMs: Refreshing Large Language Models with Search Engine Augmentation

Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong

ACL 2024 Findings

// Our dataset and method have inspired or been used for the development of Google's Gemini, Perplexity.AI's Online LLMs, You.com, and Contextual AI's RAG 2.0

The Flan Collection: Designing Data and Methods for Effective Instruction Tuning

Shayne Longpre, Le Hou, **Tu Vu**, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc Le, Barret Zoph, Jason Wei, and Adam Roberts

ICML 2023

// Google Research Blog

Mixture-of-experts meets instruction tuning: A winning combination for large language models

Sheng Shen, Le Hou, Yanqi Zhou, Nan Du, Shayne Longpre, Jason Wei, Hyung Won Chung, Barret Zoph, William Fedus, Xinyun Chen, **Tu Vu**, Yuexin Wu, Wuyang Chen, Albert Webson, Yunxuan Li, Vincent Zhao, Hongkun Yu, Kurt Keutzer, Trevor Darrell, and Denny Zhou

ICLR 2024

SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer

Tu Vu, Brian Lester, Noah Constant, Rami Al-Rfou, and Daniel Cer

ACL 2022

Overcoming Catastrophic Forgetting in Zero-Shot Cross-Lingual Generation

Tu Vu, Aditya Barua, Brian Lester, Daniel Cer, Mohit Iyyer, and Noah Constant

EMNLP 2022

STraTA: Self-Training with Task Augmentation for Better Few-shot Learning

Tu Vu, Thang Luong, Quoc Le, Grady Simon, and Mohit Iyyer

EMNLP 2021

Exploring and Predicting Transferability across NLP Tasks

Tu Vu, Tong Wang, Tsendsuren Munkhdalai, Alessandro Sordoni, Adam Trischler, Andrew Mattarella-Micke, Subhransu Maji, and Mohit Iyyer

EMNLP 2020

FUNDING & RESEARCH SUPPORT

Sentient Labs (\$35,000)	2026
PI: Tu Vu	
Google Research Award (\$50,000)	2025
PI: Tu Vu	
Google DeepMind Research Award (\$30,000)	2025
PI: Tu Vu	
Amazon - VT Faculty Research Award (\$130,000)	2025
PI: Tu Vu	
Google AI Studio Tier 1	2025
PI: Tu Vu	
VT New Faculty Mentoring Grant (\$1,500)	2025
PI: Tu Vu	
Adobe Research Gift (\$5,000)	2024
PI: Tu Vu	

ADVISING

PHD ADVISEES:

Weiyuan Chen, 1 st year PhD student	Fall 2025 — <i>present</i>
Yu-Min Tseng, 1 st year PhD student	Fall 2025 — <i>present</i>
Rishab Balasubramanian, 2 nd year PhD student	Fall 2024 — <i>present</i>
Thinh Pham, 2 nd year PhD student	Fall 2024 — <i>present</i>
Quyet Do, 2 nd year PhD student	Fall 2024 — <i>present</i>
Pin-Jie Lin, 2 nd year PhD student	Fall 2024 — <i>present</i>

MASTERS ADVISEES:

Rituraj Sharma, 1 st year MS student	Spring 2025 — <i>present</i>
Noah Provenzano, 2 nd year MS student	Spring 2025 — <i>present</i>

UNDERGRAD ADVISEES:

Mikaela Cankosyan, senior student	Spring 2026 — <i>present</i>
Nguyen Nguyen, junior student	Spring 2025

OTHERS:

Zhenting Qi, Student Researcher @ Google	Summer — Fall 2025
Prateek Yadav, Research Intern @ Google DeepMind	Summer 2024 — Spring 2025
Simeng Han, Student Researcher @ Google DeepMind	Summer 2024 — Spring 2025
Salaheddin Alzubi, Masters student @ UMass Amherst	Fall 2022 — Spring 2023
Dheeraj Mekala, PhD student @ UCSD	Spring — Summer 2022

TEACHING

CS 6804: Frontier AI Systems	Spring 2026
CS 4804: Introduction to AI	Fall 2025
CS 5624: Natural Language Processing	Spring 2025

RECENT INVITED TALKS & LECTURES

Efficient Model Development through Fine-tuning Transfer Lightning talk @ The Amazon-Virginia Tech AI Workshop	October 2025
Attention Mechanisms and Transformers Two lectures @ The New Turing Institute	August 2025
SealQA: Raising the Bar for Reasoning in Search-Augmented Language Models Open AGI Symposium @ UC Berkeley , hosted by Sentient	August 2025
Multimodal LLMs & Test-time Scaling Lecture @ The New Turing Institute	June 2025
Efficient Model Development in the Era of Large Language Models Qualcomm	November 2024
Efficient Model Development in the Era of Large Language Models Mila / McGill NLP Seminar	October 2024
Efficient Adaptation of Large Language Models Graph Neural Networks Reading Group, Google	November 2023
Effective and Efficient Transfer Learning in the Era of Large Language Models Faculty job talk	Spring 2023
Overcoming Catastrophic Forgetting in Zero-Shot Cross-Lingual Generation Parameter Efficient Tuning Methods Sync, Google	October 2022
Transfer Learning with Large-scale Language Models Lecture @ The New Turing Institute	August 2022

The Appeal of Parameter-efficient Transfer Learning
Natural Language Accelerated Team, **Google**

June 2022

SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer
Parameter Efficient Tuning Methods Sync, **Google**

December 2021

ACADEMIC SERVICE

Area Chair for ACL 2026, 2025, 2024; EMNLP 2025, 2024, NAACL 2025; COLING 2025

Session Chair (Machine Learning for NLP) @ EMNLP 2024

Program Committee/Reviewer for NATURE 2025, NEURIPS 2025, 2024; ICLR 2026, 2025; ICML 2026, 2025; TL4NLP@NEURIPS 2022; MCDC@ICLR 2025; COLM 2025, 2024; AAAI 2023; ACL 2022, 2021, 2020, 2019; EMNLP 2022, 2021; NAACL 2022, 2021; COLING 2020; CoNLL 2019; INLG 2020, 2019

SELECTED MEDIA

GEMINI: Google AI Blog	2023
FRESHLLMs: ZDNET	2023
THE FLAN COLLECTION: Google Research Blog	2023
SPoT: Headlines of Google AI's Natural Language Accelerated Newsletter	Q1, 2022

SELECTED AWARDS & HONORS

Google Student Researchships	2020 — 2023
UMass Amherst Graduate Assistantships	2016 — 2023
Honda Y-E-S Award for young engineers and scientists, Vietnam <i>// in the top 10 nationally</i>	2013
Outstanding Academic and Co-curricular Achievements, Vietnam National University	2013
Prominent Young Figure Award, Vietnam National University	2010 & 2012
First Runner-up Prize, International Programming Contest, Japan <i>// ranked 2nd among 64 teams internationally</i>	2011
Outstanding Young Talent of the Capital City, Vietnam <i>// in the top 100 most outstanding young talents selected from a wide range of fields</i>	2010
Champion Prize, National Mathematical Olympiad, Vietnam <i>// ranked 1st among more than 600 contestants nationally</i>	2010
A number of prizes in National/International Olympiads (in both Mathematics and Informatics)	2009 — 2013
A number of academic scholarships for undergraduate students	2009 — 2013

PATENTS

∗: original inventor

Frozen Model Adaptation Through Soft Prompt Transfer

Tu Vu[∗], Daniel Cer, Noah Constant, Brian Lester, Rami Al-Rfou

U.S. Patent Application, 17/863,840

Task Augmentation and Self-training for Improved Few-shot Learning

Thang Luong, **Tu Vu**[∗], Quoc Le, Grady Simon

U.S. Patent Application, 17/826,690