

TU VU

tuv@vt.edu

<https://tuvllms.github.io>

APPOINTMENTS

Virginia Tech

Assistant Professor, Computer Science

Research Interests: *large language models*, *transfer learning*

August, 2024 —

Google Research

Research Scientist

July, 2023 — *present*

EDUCATION

University of Massachusetts, Amherst

M.S/PH.D. in Computer Science

Advisor: [Mohit Iyer](#)

2016 — 2023

2018 — 2023

Vietnam National University, Hanoi

B.S. Honors Program in Computer Science

Class rank: 1st out of 100, *highest distinction*

2009 - 2013

PROFESSIONAL EXPERIENCE

Google Deepmind

Student Researcher *with* [Thang Luong](#) & [Quoc Le](#)

Fall 2022 — Spring 2023

Google Deepmind

Research Intern *with* [Thibault Sellam](#) & [Elizabeth Clark](#)

Summer 2022

Google Deepmind

Student Researcher *with* [Noah Constant](#)

Winter 2021 — Spring 2022

Google Research

Research Intern & Student Researcher *with* [Daniel Cer](#) & [Noah Constant](#)

Summer 2021 — Fall 2021

Google Deepmind

Student Researcher *with* [Thang Luong](#) & [Quoc Le](#)

Winter 2020 — Spring 2021

Google Research

Research Intern *with* [Grady Simon](#) & [Zi Yang](#)

Summer 2020

Microsoft Research

Research Intern *with* [Tong Wang](#) & [Tsendsuren Munkhdalai](#) & [Adam Trischler](#)

Summer 2019

SELECTED PREPRINTS & PUBLICATIONS

For an up-to-date list of my research papers, please see my [Google Scholar](#) profile or my [Semantic Scholar](#) profile.

1. [FreshLLMs: Refreshing Large Language Models with Search Engine Augmentation](#)
Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong
arXiv preprint 2023
2. [The Flan Collection: Designing Data and Methods for Effective Instruction Tuning](#)
Shayne Longpre, Le Hou, **Tu Vu**, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V Le, Barret Zoph, Jason Wei, and Adam Roberts
ICML 2023
3. [SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer](#)
Tu Vu, Brian Lester, Noah Constant, Rami Al-Rfou, and Daniel Cer
ACL 2022
4. [Mixture-of-experts meets instruction tuning: A winning combination for large language models](#)
Sheng Shen, Le Hou, Yanqi Zhou, Nan Du, Shayne Longpre, Jason Wei, Hyung Won Chung, Barret Zoph, William Fedus, Xinyun Chen, **Tu Vu**, Yuexin Wu, Wuyang Chen, Albert Webson, Yunxuan Li, Vincent Zhao, Hongkun Yu, Kurt Keutzer, Trevor Darrell, and Denny Zhou
arXiv preprint 2023
5. [Exploring and Predicting Transferability across NLP Tasks](#)
Tu Vu, Tong Wang, Tsendsuren Munkhdalai, Alessandro Sordoni, Adam Trischler, Andrew Mattarella-Micke, Subhransu Maji, and Mohit Iyyer
EMNLP 2020
6. [Overcoming Catastrophic Forgetting in Zero-Shot Cross-Lingual Generation](#)
Tu Vu, Aditya Barua, Brian Lester, Daniel Cer, Mohit Iyyer, and Noah Constant
EMNLP 2022
7. [STraTA: Self-Training with Task Augmentation for Better Few-shot Learning](#)
Tu Vu, Minh-Thang Luong, Quoc Le, Grady Simon, and Mohit Iyyer
EMNLP 2021